

Changepoint Analysis Using R

Robert Maidstone




17 October, 2018

What are Changepoints?

change-point *n.* (a) a point at which something changes; the point on the scale of a measuring device representing this (now *rare*); (b) *Statistics* the point at which the probability distribution of a sequence of random variables changes (frequently *attributive*, as **change point analysis**, **change point problem**, etc.).


Change in mean

```
set.seed(14)
m<-5
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
means <- rnorm(m+1,0,4)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  y[j]<-rnorm(length(j),means[i],1)
}
```




Change in mean

```
set.seed(14)
m<-5
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
means <- rnorm(m+1,0,4)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  y[j]<-rnorm(length(j),means[i],1)
}
```




Change in variance

```
set.seed(12)
m<-5
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
sd <- runif(m+1,1,20)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  y[j]<-rnorm(length(j),0,sd[i])
}
```




Change in variance

```
set.seed(12)
m<-5
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
sd <- runif(m+1,1,20)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  y[j]<-rnorm(length(j),0,sd[i])
}
```




Change in trend

```
set.seed(110)
m<-3
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
slope <- rnorm(m+1,0,.01)
intercept <- rnorm(1,0,1)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  if(i==1){
    for(jind in j){
      y[jind]<-intercept+(jind-true_cps[i])*
        slope[i] + rnorm(1,0,1)
    }
  }else{
    for(jind in j){
      y[jind]<-y[j[1]-1]+(jind-true_cps[i])*
        slope[i] + rnorm(1,0,1)
    }
  }
}
```



Change in trend

```
set.seed(110)
m<-3
n<-1000
true_cps <- c(0,sort(sample(1:(n-1),m)),n)
slope <- rnorm(m+1,0,.01)
intercept <- rnorm(1,0,1)
y<-c()
for(i in 1:(m+1)){
  j <- (true_cps[i]+1):true_cps[i+1]
  if(i==1){
    for(jind in j){
      y[jind]<-intercept+(jind-true_cps[i])*
        slope[i] + rnorm(1,0,1)
    }
  }else{
    for(jind in j){
      y[jind]<-y[j[1]-1]+(jind-true_cps[i])*
        slope[i] + rnorm(1,0,1)
    }
  }
}
```



Real World Examples

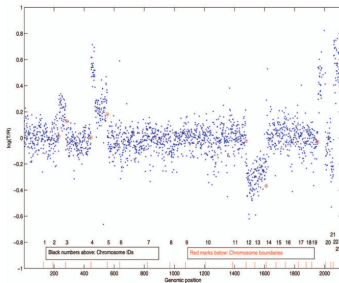
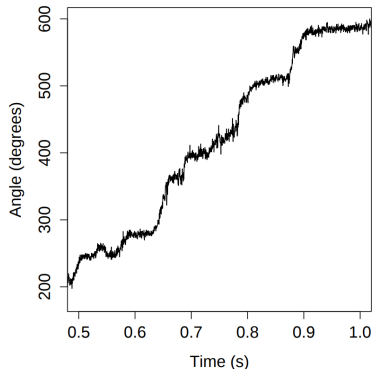


Fig. 5.
Genome of the breast tumor S1514 [39].

(a) Copy number at genomic positions in a human breast tumor sample (Chen & Wang, 2009).



(b) Rotation of a bacterial flagella motor (Maidstone, 2016).

Many changepoint methods

Many different changepoint algorithms exist

- Exhaustive Search
- Optimal Partitioning
- PELT
- FPOP (and R-FPOP)
- CROPS
- Segment
Neighbourhood
Search
- pDPA
- SNIP
- Binary Segmentation
- WBS
- CBS
- SMUCE
- SMOP
- ED-PELT
- E-Divisive
- ECP

Binary Segmentation

Input: A set of data of the form, (y_1, y_2, \dots, y_n) where $y_i \in \mathbb{R}$.
A test statistic on the data $\Gamma(\cdot)$,
An estimator of changepoint location $\hat{\tau}(\cdot)$,
A rejection threshold c .

Initialise: Let $\mathcal{CP} = \emptyset$ and $\mathcal{S} = \{[1, n]\}$;

Iterate: while $\mathcal{S} \neq \emptyset$ do

- Choose an element of \mathcal{S} ; denote this element as $[s, t]$;
- if $\Gamma(\mathbf{y}_{s:t}) < c$ then
 - remove $[s, t]$ from \mathcal{S}
- if $\Gamma(\mathbf{y}_{s:t}) \geq c$ then
 - remove $[s, t]$ from \mathcal{S} ;
 - calculate $r = \hat{\tau}(\mathbf{y}_{s:t}) + s - 1$, and add r to \mathcal{CP} ;
 - if $r \neq s$ add $[s, r]$ to \mathcal{S} ;
 - if $r \neq t - 1$ add $[r + 1, t]$ to \mathcal{S} ;

Output: The changepoints recorded in \mathcal{CP} .

Binary Segmentation using “wbs”

Baranowski and Fryzlewicz (2015)

```
library(wbs)




sbs(y) -> sbs.out
s.cpt <- changepoints(sbs.out,th =4)

s.cpt$cpt.th[[1]] %>% sort
## [1] 264 637 955 979

wbs(y,
     M=5000) -> wbs.out

w.cpt <- changepoints(wbs.out,th=4)

w.cpt$cpt.th[[1]] %>% sort
## [1] 264 521 637 955 979
```



Cost function representation

Most changepoint detection methods boil down to minimising the sum of some cost function, $\mathcal{C}(\cdot)$, over the segments.

$$\min_{\tau_{1:m}, m} \left[\sum_{j=0}^{m+1} \mathcal{C}(\mathbf{y}_{\tau_{j+1}:\tau_{j+1}}) \right]$$

This cost function could be a number of things:

- 1 Negative log-likelihood,
- 2 Negative posterior,
- 3 Minimum Description Length.

Dynamic Programming Methods

Optimal Partitioning: Optimisation based sum of optimal up to last changepoint and the cost between last changepoint and current time (plus a penalty to avoid over fitting).

$$F(\tau^*) = \min_{0 \leq \tau < \tau^*} [F(\tau) + \mathcal{C}(\mathbf{y}_{(\tau+1):\tau^*}) + \beta].$$

Segment Neighbourhood Search: Optimisation for k segments based on optimal for $k - 1$ segments plus cost for new segment.



$$q_{1,j}^k = \min_{v \in \{1, \dots, j-1\}} [q_{1,v}^{k-1} + q_{v+1,j}^1].$$

The changepoint package

- Functions for changepoint analysis.
- Can use various changepoint detection methods:
 - Binary Segmentation
 - AMOC
 - Segment Neighbourhood
 - PELT (Optimal Partitioning)
- With various choices of penalty function/value.
- And either a Gaussian or CUSUM test statistic.
- Three *headliner* functions:
 - `cpt.mean`
 - `cpt.var`
 - `cpt.meanvar`
- Authors: Rebecca Killick, Kaylea Haynes, Idris Eckley, Paul Fearnhead, Jamie Lee.



Change in mean using “changepoint”

```
changepoint::cpt.mean(y_mean,  
                      penalty = "BIC",  
                      method = "PELT"  
                      ) -> cpt_object #S4  
  
cpt_object %>% changepoint::likelihood()  
  
##      -2*logLik -2*Loglike+pen  
##      2839.115      2894.377  
  
cpt_object %>% changepoint::ncpts()  
  
## [1] 4  
  
cpt_object %>% changepoint::cptests() -> est_cps  
qpcR:::rbind.na(tru_cps,est_cps)  
  
##      [,1] [,2] [,3] [,4] [,5]  
## tru_cps 254 551 637 955 979  
## est_cps 264 637 955 979 NA
```



Change in variance using “changepoint”





```
changepoint::cpt.var(y_var,  
                     penalty = "BIC",  
                     method = "PELT"  
                     ) -> cpt_object #S4  
  
cpt_object %>% changepoint::likelihood()  
  
##      -2*logLik -2*Loglike+pen  
##      4395.668      4464.746  
  
cpt_object %>% changepoint::ncpts()  
  
## [1] 5  
  
cpt_object %>% changepoint::cpts() -> est_cps  
qpcR::rbind.na(tru_cps,est_cps)  
  
##      [,1] [,2] [,3] [,4] [,5]  
## tru_cps  70  169  269  817  940  
## est_cps  70  169  268  815  940
```



Optimal Partitioning: Optimisation based sum of optimal up to last changepoint and the cost between last changepoint and current time (plus a penalty to avoid over fitting).

$$F(\tau^*) = \min_{0 \leq \tau < \tau^*} [F(\tau) + \mathcal{C}(\mathbf{y}_{(\tau+1):\tau^*}) + \beta].$$

Penalty Choice




CROPS

```
cpt.mean(y,penalty = "CROPS",  
         pen.value=c(5,1000),  
         method="PELT",  
         class=FALSE) -> crops.out
```

```
## [1] "Maximum number of runs of algorithm = 19"  
## [1] "Completed runs = 2"  
## [1] "Completed runs = 3"  
## [1] "Completed runs = 5"  
## [1] "Completed runs = 7"  
## [1] "Completed runs = 11"  
## [1] "Completed runs = 16"
```



```
crops.out$cpt.out %>% .[,1:5]
```

	[,1]	[,2]	[,3]	[,4]	[,5]
## beta_interval	5.0000	5.199941	5.642361	6.626451	8.256657
## numberofchangepoints	20.0000	18.000000	16.000000	12.000000	11.000000
## penalised_cost	890.7315	901.131370	912.416092	938.921896	947.178554




Change in trend with Prophet

```
df <- data.frame(  
  ds=lubridate::as_date("16/10/18")+  
    (1:length(y)),  
  y= y)  
  
m <- prophet(df)  
  
forecast <- predict(m)  
  
plot(m, forecast)+add_changepoints_to_plot(m)  
  
prophet_plot_components(m, forecast)
```




Peyton Manning (Prophet example)

```
df<-read.csv(  
  "~/changepointsinR/  
  example_wp_log_peyton_manning.csv")  
m <- prophet(df)  
  
future <- make_future_dataframe(  
  m,  
  periods = 365)  
  
forecast <- predict(m, future)  
  
plot(m, forecast)+add_changepoints_to_plot(m)
```






Peyton Manning (Prophet example)



Peyton Manning (Prophet example)

```
prophet_plot_components(m, forecast)
```



- Many changepoint methods available
 - Not discussed; Multivariate, online changepoint detection, non parametric (`changepoint.np`, `ecp`) and more.
- `changepoint` and `wbs` packages are a good place to start.
- Computational time can be a major issue in changepoint detection, meaning bespoke solutions can be best.
- Code not always in the best form, but big forecasting packages such as `prophet` use changepoints too.